

Shafeeq Ahamed Saiffudin

US Permanent Resident | San Francisco, CA | +1 (925) 319-8712 | shafeeq.jms@gmail.com | [LinkedIn](#) | [GitHub](#)

PROFESSIONAL SUMMARY

DevOps Engineer with 6+ years building cloud infrastructure and running production systems at scale. Saved \$834K through AWS optimization and cloud migrations. Specializing in model-serving infrastructure (GPU), Terraform, Ansible, and Kubernetes.

TECHNICAL SKILLS

Cloud & Infrastructure: AWS (EC2, ECS, EKS, S3, Batch, Bedrock, CloudWatch, Secrets Manager, VPC, Auto Scaling, CodePipeline), Azure, GCP, Cloudflare

Containers & Orchestration: Docker, Kubernetes, Helm, Amazon EKS, AWS ECS

Infrastructure as Code & CI/CD: Terraform, Ansible, CloudFormation, GitHub Actions, AWS CodePipeline, Git

Monitoring & Observability: Datadog, Grafana, Prometheus, PagerDuty, Icinga2, Mezmio, CloudWatch

Scripting & Automation: Python, Bash, Shell scripting

Security & Compliance: NIST 800-53, Nessus, OWASP ZAP, Probely, Okta OIDC, AWS Secrets Manager

Databases & Storage: MySQL, PostgreSQL, Elasticsearch, Redis

Networking: HAProxy, AWS ALB/NLB, DNS, Route53, VPC, Security Groups

AI/ML Infrastructure: GPU cluster management (Nvidia H100), multi-node training, ML model deployment pipelines, AI-assisted development (Claude Code, Cursor), AWS Bedrock, OCR pipelines

Operating Systems: Linux (Ubuntu, Amazon Linux)

CERTIFICATIONS

AWS CERTIFIED - DEVOPS ENGINEER PROFESSIONAL

AWS CERTIFIED - SYSOPS ADMINISTRATOR ASSOCIATE

EXPERIENCE

ColorKloud

San Francisco, CA, USA

Founding DevOps Consultant | Sep 2024 – Present

Founded consulting practice specializing in cloud infrastructure and AI-powered serverless applications for startups and small businesses.

- Deployed 17 production serverless applications (observability tools, fintech platforms, ecommerce, operations dashboards) on Cloudflare Workers with D1 database, integrating Stripe, RevenueCat, Twilio, Postmark, Supabase, Clerk, and Okta
- Built AI-powered applications using Cloudflare Workers AI bindings and AWS Bedrock for OCR pipelines running on ECS
- Cut client hosting costs by 40%, measured by reducing monthly spend from \$500 to \$300 per client, by migrating on-premises apps to the cloud with Terraform and Ansible
- Sped up release cadence, measured by going from weekly to daily deployments, by using Claude Code and Cursor to write and validate infrastructure code
- Reduced feature delivery time, measured by shipping bug fixes in hours instead of days, by leveraging AI-assisted development tools

Atmo Inc

San Francisco, CA, USA

Platform Engineer (via ColorKloud B2B contract) | Mar 2025 – Nov 2025

Built and maintained GPU-based infrastructure for ML weather forecasting models serving the Department of Defense and government entities.

- Architected AWS ECS infrastructure with Nvidia H100 GPUs scaled to support 70 concurrent multi-node weather forecast jobs for DoD clients, improving deployment speed by 40% and scaling efficiency by 2x
- Cut deployment time by 83%, measured by reducing deploys from 1 hour to 10 minutes, by writing Dockerfiles and automating containerized ML model deployments

- Wrote Python monitoring scripts, measured by S3 file freshness checks and automated alerts, by integrating with Datadog to detect missing forecasts and trigger backups
- Saved \$750K annually, measured by AWS billing reduction, by rightsizing H100 GPU instances, implementing AWS Batch scheduling, and optimizing S3 lifecycle policies
- Shipped frontend releases 2-3x daily, measured by going from weekly to multiple daily deployments in under 20 minutes, by building CI/CD with GitHub Actions and CloudFront
- Kept production at 99.9% uptime, measured by uptime monitoring across DoD systems, by running 24x7 on-call rotation and proactive incident mitigation
- Passed NIST 800-53 certification, measured by successful DoD audit, by conducting security audits and compliance evaluations

LocationIQ - Unwired Labs India Pvt Ltd

Hyderabad, India

DevOps Engineer & SRE | Mar 2020 – Mar 2024

- Saved \$84K annually, measured by infrastructure cost reduction, by migrating 40% of on-premise workloads to AWS with capacity optimization
- Dropped p95 response time by 75%, measured by reducing latency from 100-200ms to under 50ms, by building HAProxy config to route traffic based on latency
- Reduced critical vulnerabilities by 80%, measured by dropping findings from 10-15 per quarter to 2-3 per quarter, by adding security pre-checks to production rollout process
- Automated CI/CD pipelines, measured by deployment consistency and speed, by implementing AWS CodePipeline for continuous integration and delivery
- Ran quarterly security scans, measured by finding and fixing critical vulnerabilities within 48 hours, by using Nessus, OWASP ZAP, and Probely and coordinating fixes with developers
- Eliminated downtime during OS upgrades, measured by 100% success rate across 40+ servers, by executing zero-downtime Ubuntu upgrade process
- Maintained monitoring with Grafana, Mezmo, PagerDuty, and Icinga2, measured by detecting and resolving app and security incidents, as first on-call responder
- Kept servers consistent across AWS, bare metal, and colocation, measured by preventing config drift, by managing infrastructure with Ansible

Zudioz

Chennai, India

DevOps Engineer | Oct 2018 – May 2020

- Built 12 e-commerce sites, measured by sites deployed to production, by developing on Joomla, Magento, and WordPress for manufacturing clients
- Improved disaster recovery MTTR by 94%, measured by reducing database failure recovery from 2 days to 3 hours, by implementing MySQL on-site and off-site backup automation
- Ran MySQL replication, measured by maintaining consistency across 27 servers, by setting up master-replica configuration
- Maintained AWS EC2 infrastructure, measured by uptime and availability, for internal tools and client applications

EDUCATION

Anna University, Chennai, India

B.E. in Electrical & Electronics Engineering | Aug 2012 - Aug 2016

- *Foundation in embedded systems, control systems, and microcontroller programming.*

KEY ACHIEVEMENTS

GPU/ML Infrastructure: Architected Nvidia H100 GPU infrastructure scaled to 70 concurrent multi-node forecast jobs for DoD weather forecasting at 99.9% uptime

Cost Optimization: \$834K total infrastructure savings (\$750K at Atmo, \$84K at LocationIQ)

AI/Serverless: Built 17 production AI-powered serverless apps with Cloudflare Workers AI, AWS Bedrock, Stripe, RevenueCat integrations

Automation Impact: 83% deployment time reduction, 2-3x daily release frequency, 94% disaster recovery improvement

Security & Compliance: NIST 800-53 DoD certification, 80% reduction in critical vulnerabilities

Agent-Driven Development: Proficient in Claude Code, Pi harness, and Cursor for accelerated infrastructure development